

What 16,199 Discretionary Land-Use Decisions Look Like: San Francisco Planning Commission, 1998–2026

Dan Post

September 2026

Abstract

This memo reports what the item-level dataset contains now that it exists. Every item the San Francisco Planning Commission heard between 1998 and mid-2026 that carries a case number — **16,199 items across 8,987 distinct cases** — has been extracted into 29 structured fields. It describes the composition of the docket, what the Commission does to what comes before it, how long projects take, how individual commissioners vote, where in the city the activity is, and which parts of the Planning Code are actually invoked.

Four findings organise the rest. **Outright denial is close to non-existent** — 1.4% of items across twenty-eight years, and effectively zero since 2020 — while **conditioning has more than tripled**, from 13.8% of items in 1998 to 45.7% in 2023. The Commission has not become more permissive; it has moved from a denial technology to a conditioning technology. **Delay is the second margin**: 45% of all items are a repeat appearance of a case already heard, and the 90th-percentile case now takes a year from first hearing to last, up from five months in the late 1990s. **Discretionary review has collapsed** as a share of the docket, from 28.6% to 8.6%, following the 2013 reforms.

The memo closes on what the data cannot yet do. The regulatory burden a developer faces is not a disposition but an *expectation* formed before filing — over delay, over conditions, over the probability of either. The dataset now measures realised outcomes well. Turning those into a measure of *ex ante* cost and uncertainty is the modelling problem the paper rests on, and §9 sets out four routes into it. A tested linkage to the city’s building-permit records — **97.6% of all 2,295 parsed permit numbers match** — shows that construction outcomes are observable on the other side of the entitlement.

Contents

1	The data	3
1.1	What is in it	3
1.2	Two coverage gaps to keep in view	3
2	The composition of the docket	3
3	What the Commission does	5
3.1	Denial is not the instrument	5
3.2	Conditioning is	5
3.3	Continuance is the third margin	6
3.4	The disposition depends on what is asked	6
4	Delay	6

5	Commissioners	6
6	Geography	9
7	Which parts of the Planning Code are actually used	10
8	What the patterns say	11
8.1	Stringency moved from the extensive to the intensive margin	11
8.2	The binding constraint is time, and specifically its variance	11
8.3	Discretionary review's collapse is a natural experiment	11
9	From this dataset to a cost of uncertainty	11
9.1	The problem with <code>conditions_imposed</code> as it stands	11
9.2	Route one: realised delay as a predictable quantity	12
9.3	Route two: the staff recommendation as a prior	12
9.4	Route three: speakers as a mobilisation shock	12
9.5	Route four: the option value of not filing	12
10	Linking to the city's permit records	13
10.1	Permit numbers parse and they match	13
10.2	What the match buys	13
10.3	The parcel route is broader but noisier	13
10.4	What is not yet established	14
11	Where the data lives	14
12	Limitations	14

1 The data

1.1 What is in it

One row is one item on one agenda: a request heard by the Commission on a date. The corpus runs from January 1998 to the last scraped meeting of 2026 and holds **16,199 items** carrying a case number, drawn from 985 meeting documents. Blocks with no case number — public comment, the Director’s report, the weekly recap, adjournment — are agenda scaffolding and are excluded by rule.

Extraction is Claude Haiku 4.5 with six retrieved, era-matched worked examples, run over the Batch API against a schema of 29 fields. On a frozen 139-item test half that no prompt was tuned against, the configuration scores **96.8%** at the field level: 97.4% on the 1998–2014 HTML minutes and 94.0% on the 2015–2026 PDF minutes. The companion memo, *Item-Level Field Extraction*, documents the method, the metric and the failure modes; this memo takes the output as given and asks what it says.

1.2 Two coverage gaps to keep in view

1998–2000 is thin. Those years hold 12 meeting documents each against roughly 40 in every year from 2001. The Commission met weekly throughout; the archive does not. Rates computed within those years are usable, counts are not comparable, and the composition figures should be read as describing the meetings that survive.

2025–2026 is incomplete. Eight documents each, against 34–41 in the preceding years, because the scrape reached the archive before the archive reached the present. Every count for those years is a floor, and the memo’s figures mark the boundary rather than letting a reader infer a collapse in Commission activity that did not happen.

2 The composition of the docket

Table 1 gives the numbers. Three movements matter.

Discretionary review collapsed. DR was 28.6% of the docket in 1998 and 38.8% in 2003 — at its peak, the single largest category, and the mechanism by which any member of the public could pull an as-of-right building permit in front of the Commission. By 2013 it was 12.5% and by 2025, 8.6%. The timing points at the 2010–2013 DR reforms, which raised the bar for a request to be heard at all. This is the clearest structural break in the dataset and it changes what the Commission’s docket *is*: a body that spent a third of its agenda reviewing individual permits now spends a twelfth.

Conditional use grew to fill the space, from 35.9% to 48.4% at its 2023 peak. Some of this is composition rather than volume — the denominator shrank as DR left — but the absolute count of CU items is also higher in the 2020s than the 2000s despite a smaller docket overall.

Legislative business roughly doubled, from 6.6% to 15.6%, and *informational* items went from a rounding error to 14.8% of the 2025 docket. The Commission increasingly spends its time on rules rather than on projects. For a study of regulation as a market, that is not a nuisance category: it is the supply side of rule-making appearing in the same series as the demand side of individual permissions.




Figure 1: Items heard per year (top) and their composition (bottom). The thin years at each end are archive coverage, not Commission activity.




Figure 2: Outcome rates. Faint lines are the annual series; bold lines are a three-year centred mean.

3 What the Commission does

Table 2 is the central result of this memo.

3.1 Denial is not the instrument

Across twenty-eight years the Commission disapproved **1.4%** of the items before it. In 1998 the rate was 4.1%; since 2020 it rounds to zero. If regulatory stringency were measured by refusal, San Francisco would look like one of the most permissive jurisdictions in the country, which is not a conclusion anybody believes.

3.2 Conditioning is

Over the same period the share of items carrying conditions rose from **13.8%** to **45.7%**. The Commission has substituted along a margin that a denial-rate measure cannot see: it says yes, and attaches terms. Approval with conditions is the modal outcome of a modern conditional-use hearing.

This substitution is the single most important pattern for the paper’s framing. A discretionary tax that operates through conditions rather than refusals is invisible to any measure of approval probability, and it is precisely the kind of burden that is costly to anticipate — which is where uncertainty enters.

3.3 Continuance is the third margin

39% of all items end in a continuance, though the rate has fallen from 46% to 25%. Table 2 understates how much of the docket this represents, because a continued item returns: **45% of all 16,199 items are a repeat appearance of a case already heard**. Nearly half the Commission’s agenda is re-hearing.

3.4 The disposition depends on what is asked

Pooling all request types hides that these are different proceedings. Conditional use is approved 54.9% of the time and continued 35.9%. Discretionary review is approved 1.3% of the time — because “approval” is not its vocabulary: 26.8% end in *did not take DR* and 15.4% in *took DR and approved*, so the modal DR outcome is the permit proceeding essentially untouched. Variances are continued 44.6% of the time and approved 15.4%, which reflects that the Zoning Administrator, not the Commission, is the usual variance authority.

4 Delay

A case heard more than once is a project the Commission saw and did not finish with. Chaining items by case number gives 8,987 cases, of which **37.9% were heard more than once**; the record is a case that took 3,423 days — nine years and four months — across its appearances.

Table 3 shows the distribution by cohort. The median has been stable at about five weeks, which is roughly the Commission’s own calendar cycle: a routine continuance moves an item two or three meetings. **The tail is what moved.** The 90th percentile went from 147 days for cases first heard in the late 1990s to **360 days** for the 2015–2019 cohort. The typical case is not slower. The unlucky case is more than twice as slow.

That is a statement about *variance*, and it matters more than the median for a developer deciding whether to file. Elapsed time also sorts on outcome in a way that is intuitive but worth having measured: cases ending in approval run a median 28 days from first to last hearing, those ending *continued indefinitely* run 84, and those ending in no action run 116.

Two cautions. First, this measures hearing-to-hearing time, not the full entitlement clock: it begins when an item first reaches an agenda, which is already well downstream of filing. Second, recent cohorts are right-censored — a case first heard in 2024 cannot yet show a three-year tail — so the recent decline in Table 3 is partly an artefact and is flagged as such rather than read as an improvement.

5 Commissioners

The minutes record roll-call votes on both sides, so this is a census rather than a sample: **79,855 individual votes by 84 named commissioners**. Table 4 gives the twenty most active.

The striking feature is how **rarely anyone dissents**. Rates run from 1.1% to 7.9%, and the Commission is close to unanimous on the overwhelming majority of items. The spread is nonetheless real and persistent: Moore (7,608 votes, 7.2%), Sugaya (7.9%) and Olague (7.8%) dissent five to seven times as often as Diamond (1.1%), Borden (1.2%) or Koppel (1.4%), and these are not small samples.

Two things this enables. First, a commissioner fixed effect is estimable for the two dozen members with enough votes, which makes it possible to ask whether *who is sitting* moves outcomes




Figure 3: Repeat hearings and elapsed time, by year of first hearing.




Figure 4: Dissent rate for commissioners with 200 or more recorded votes.

conditional on what is asked. Second, because the roll call is per item and the composition of the Commission turns over, the panel has within-project variation in composition for cases heard repeatedly — the same project, a different room.

6 Geography




Figure 5: Assessor blocks appearing most often before the Commission, 1998–2026.

84.1% of items carry a usable parcel key — an assessor block and at least one lot — covering 2,557 distinct blocks and 8,295 distinct parcels. This is the join that makes everything downstream possible, and it is the reason `project_address` was demoted to a validation field: the block-and-lot pair is what the city’s own records key on, and it scores 98–100% where the address needs a geocoder.

Concentration is high. The most-heard block (3736, South of Market at the Rincon Hill/Transbay edge) appears 95 times; the top ten are almost entirely SoMa and the Downtown core — 3704, 3721, 3753, 3787, 3788 — with two North Beach blocks (131, 287) as the exception. These are the blocks where large projects, phased approvals and repeated returns concentrate.

Hot spots are a genuine finding but they need care: a block appears often both because a lot

of building is proposed there *and* because a single large project can generate a dozen items across a decade. Separating those is a matter of collapsing to the case level and is straightforward with the chain construction of §4.

7 Which parts of the Planning Code are actually used

Table 5 lists the sections cited in the request text. **1,347 distinct sections** appear at least once; the distribution is extremely skewed.




Figure 6: The twenty Planning and Administrative Code sections cited most often in the request text. The long tail is the point: 1,347 sections appear at least once, but twenty of them account for the bulk of the citations.

The result is a validity check as much as a finding. Nothing told the extractor what the Planning Code contains, yet the mapping from section to request type recovers the Code’s own structure: §303 (conditional use procedure) is cited by 2,468 conditional-use items; §309 (permit review in C-3 districts) by 252 downtown-project items; §329 (large project authorization) by 179 LPA items; §321 (office development annual limit) by 209 office-allocation items. Where the extraction assigns a request type, the cited section independently agrees.

Substantively, four sections carry most of the discretionary weight:

- **§31.04** (2,899 items) is the Administrative Code’s approval-action provision for CEQA. Its ubiquity is procedural boilerplate, but it dates the moment at which the entitlement becomes appealable.

- **§303** (2,701) is the conditional-use standard — the “necessary or desirable” test, which is the discretion.
- **§101.1** (923) is the eight Priority Policies. It appears on legislative items and is the general-welfare hook that makes almost any objection cognisable.
- **§134** (912, rear yard) and **§140** (358, dwelling-unit exposure) are the two bulk standards that generate the most variance requests — the physical constraints that bind hardest on infill.

§317, the demolition control, appears on 660 items and is the single most direct instrument of housing-stock protection in the dataset.

8 What the patterns say

8.1 Stringency moved from the extensive to the intensive margin

The headline pattern is a substitution: denial fell to zero while conditioning tripled. A regulator that rarely refuses but usually attaches terms is not a permissive regulator; it is one whose burden is expressed in the content of the approval rather than in its existence. Any measure of regulatory stringency built on approval rates will read San Francisco as permissive and will be wrong.

8.2 The binding constraint is time, and specifically its variance

Half the docket is re-hearing. The median case moves at the speed of the meeting calendar, but the 90th percentile more than doubled. For a developer with financing costs, the relevant quantity is not expected delay but the width of its distribution — and that width has grown even as the centre has not.

8.3 Discretionary review’s collapse is a natural experiment

DR fell from 38.8% of the docket to 8.6% following reforms that changed who could trigger a hearing. That is a sharp, dated, policy-induced change in the *intensity* of discretionary review — exactly the parameter the supply model needs to move — and it is identified by something other than the demand for housing.

9 From this dataset to a cost of uncertainty

The honest position is that this dataset measures **realised outcomes** well and **ex ante burden** not at all. A developer choosing whether to file faces a distribution over delay, over conditions and over approval; the dataset records one draw from it. Four routes seem worth trying, in increasing order of ambition.

9.1 The problem with `conditions_imposed` as it stands

A concrete obstacle first, because it is the binding one. `conditions_imposed` is populated on 25.9% of items — but its median length is **three characters**. It is a yes/no flag. The minutes very often say that conditions were imposed without enumerating them, and the schema faithfully records that. `modifications` does carry text (median 108 characters) but is populated on only 10.7% of items.

So the dataset can currently answer *whether* a project was conditioned, and almost never *how*. That is the gap between the pattern in §3 and a cost measure, and closing it is the highest-value next step. Three options:

1. **Motions and resolutions.** The Commission’s conditions live in the numbered motion or resolution, not in the minutes. `action_instrument` and `action_instrument_no` are extracted on 40% of items and give the document number directly. Those documents are public and are the actual text of the burden.
2. **The Planning Department’s own condition-of-approval templates.** Conditions are substantially standardised. A modest taxonomy — affordability, transportation demand management, open space, hours of operation, design review, monitoring — would convert most condition text into a handful of indicators.
3. **Cost-relevant coding.** Not every condition is costly. The distinction that matters is between conditions that change the project (units, height, envelope) and those that change its obligations (fees, monitoring, reporting). The `project_modified` flag already separates these in principle.

9.2 Route one: realised delay as a predictable quantity

The chain construction gives, for every case, the number of hearings and elapsed days. Regress that on what was observable *at first hearing* — request type, district, block, staff recommendation, speaker turnout — and the fitted value is an expected delay conditional on observables, while the residual variance is the part a developer cannot forecast. That residual is the natural first proxy for uncertainty, and it is available now without new data.

9.3 Route two: the staff recommendation as a prior

`preliminary_recommendation_category` is populated on 78.7% of items and is issued *before* the hearing. It is a public, dated signal of the expected outcome. Two quantities follow: the follow rate (how often the Commission does what staff asked), which prices how much residual discretion the Commission exercises; and the subset where the Commission departs from staff, which is where the unpredictable burden concentrates. A model in which the developer’s prior is the staff recommendation and the Commission’s deviation is the shock is both estimable and economically legible.

9.4 Route three: speakers as a mobilisation shock

Public turnout is recorded per item with a stance where the minutes mark one. Turnout is plausibly the most visible predictor of trouble, and it is observable to the developer only imperfectly at filing. Its usefulness is limited by the field’s own accuracy — 69.6% on the test half, the weakest field in the schema — so this route needs the stance work finished before it can carry weight.

9.5 Route four: the option value of not filing

The deepest version of the question is not measurable in this dataset at all, and it should be said plainly. Projects that were never filed because the process was too uncertain do not appear in Commission minutes. Any estimate of the cost of uncertainty built only on the observed docket is conditioned on selection into it. The permit records of §9 are the only route to the counterfactual:

parcels where a permit was filed and abandoned, or where an entitlement was granted and never built, are the observable trace of a project that lost its bet.

10 Linking to the city’s permit records

The dataset’s value multiplies if the entitlement can be followed to construction. This section reports a first, tested pass.

10.1 Permit numbers parse and they match

4,083 items mention a building, site or demolition permit in their request text. A permit number can be parsed from **3,588** of them, yielding 2,295 distinct permits in four printed forms. Tested against *every* one of those numbers in DataSF’s Building Permits dataset (1,295,048 rows, `i98e-djp9`), **2,239 matched — 97.6%**. The join works, and it works on the exact identifier the minutes print.

The fourth printed form is a correction to an earlier reading of this section. From about 2017 the minutes print the permit as 2019.0618.3775 — year, MMDD, serial — and a parser that insists on the older four-group 2013.10.21.9832 sees nothing. That, rather than any change in what the Commission hears, is why the permit route appeared to die after 2019. The permit-linkage memo works this through.

10.2 What the match buys

The DBI record carries `filed_date`, `issued_date`, `first_construction_document_date`, `status`, `estimated_cost`, `revised_cost`, `proposed_units` and `existing_units`. So for a matched item we can observe, on the far side of the discretionary hearing:

- **whether the permit was ever issued**, and how long after entitlement;
- **whether construction started**, via the first construction document;
- **whether it was withdrawn or expired** — the abandonment outcome that §8.5 identifies as the trace of a lost bet;
- **how much it cost** and **how many units** resulted, against how many were proposed at the hearing.

Sampled records show all four states occurring: permits complete, permits issued but not completed, and permits withdrawn after filing.

10.3 The parcel route is broader but noisier

Where no permit number is printed, the block-and-lot key reaches further, but it does not point at one permit: the median matched parcel carries a decade of plumbing, electrical and small-alteration filings irrelevant to the entitlement. Scored against the permit-number subset — which is ground truth, since there we know which permit the hearing was about — a rule that keeps only new construction, alterations and demolitions filed within three years before and one year after the hearing recovers **92% of the right permits at 0.70 precision**, narrowing a median-ten parcel history to a median-two candidate set. The permit-linkage memo reports the full trade-off curve.

10.4 What is not yet established

Whether the *conditions* attached at the hearing are observable in the permit record is now resolved, and the answer is no: across all 1,295,048 DBI rows, 87 — 0.007% — mention a condition of approval or a Commission motion at all. DBI records the permit’s own status and scope, not the Planning Commission’s conditions of approval. Recovering those means the motions and resolutions of §8.1, not DataSF. The permit join gives outcomes; it does not give the burden.

11 Where the data lives

The corpus pass is stored under `$MFHR_DATA_ROOT/extraction/corpus_v2_g3/` in three forms kept separately and never conflated:

- `raw/` — the model’s replies exactly as returned, never rewritten;
- `interim/` — unwrapped from the evidence envelope and type-coerced;
- `clean/` — normalised; the analysis input for this memo;
- `evidence/` — the supporting spans, and `verification_failures.csv`;
- `manifest.json` — batch ids, run ids, token counts, cost, prompt SHA, schema and gold versions.

The separation is not fastidiousness. Every serious defect found while building the gold set was a *normalisation* defect — one silently emptied a field on 122 of 232 gold records — and each was recoverable only because the untouched replies had been kept. Re-deriving `clean` from `raw` is a local loop; re-querying the corpus is \$69.

12 Limitations

1. **Coverage is uneven at both ends.** 1998–2000 and 2025–2026 hold roughly a fifth of the meeting documents of the intervening years. Rates are usable; counts are not.
2. **Field accuracy is 96.8%, not 100%.** At the item level the errors compound: an analysis conditioning on five fields at once is working with materially fewer clean rows than the field-level rate suggests, and whole-record accuracy should be reported wherever that is the operative unit.
3. **conditions_imposed is a flag, not a description,** which is the binding limitation on measuring stringency and the subject of §8.1.
4. **Case-number chains are an imperfect project identifier.** A project can change case number across a re-filing, and one case number can cover several agenda sub-items heard the same day. The chains in §4 are a lower bound on how often projects return.
5. **Delay is measured from first hearing,** which is already downstream of filing. The pre-hearing clock — often the longer one — is invisible here and lives in the Department’s own records.
6. **Speaker stance is the weakest field** (69.6%) and the stance rule changed after the scored run. Any mobilisation analysis should wait for a re-run.

7. **No causal claim is made anywhere in this memo.** The DR collapse in §7.3 is described as a candidate natural experiment, not as an estimate.

Table 1: Composition of the Commission’s item-level docket, per cent of items heard that year. Families group the 17 `request_type` values; ‘Legislative’ is Planning Code, map and General Plan amendments.

Year	Conditional use	Discretionary review	Variance	Legislative	Large project /downtown	CEQA /appeals	Historic	Informational	Other	Items
1998	35.9	28.6	3.9	6.6	5.6	14.9	2.9	0.7	1.0	715
2003	33.8	38.8	9.1	6.8	2.1	6.4	1.0	0.8	1.4	809
2008	27.1	24.9	8.5	17.7	5.6	8.0	1.4	4.3	2.5	851
2013	42.3	12.5	8.0	15.7	8.0	7.4	0.2	4.1	1.8	513
2018	37.6	19.6	7.6	15.6	5.0	4.5	0.4	5.2	4.5	540
2023	48.4	14.4	2.9	15.6	5.8	3.1	0.9	4.4	4.4	455

Table 2: Commission outcomes, per cent of items heard that year. ‘Conditions’ and ‘modified’ are the orthogonal flags, not dispositions: a project can be approved with conditions and modified at once.

Year	Continued	Conditions	Modified	Disapproved	Took DR
1998	46.4	13.8	9.2	4.1	5.4
2003	48.2	27.1	11.2	0.6	5.5
2008	44.5	20.7	8.6	1.7	5.0
2013	27.2	40.0	7.4	1.6	3.0
2018	35.9	29.7	9.0	1.5	5.0
2023	24.6	45.7	7.5	0.0	3.8
All years	39.0	26.4	8.8	1.4	5.0

Table 3: Elapsed time from first to last hearing, for cases heard more than once, by five-year cohort of first hearing. A case is a `case_number`; the corpus ends 2026-03-26, so recent cohorts are right-censored and their tails are understated.

First heard	Cases	Median days	90th pct	Max
1998–1999	314	35	147	665
2000–2004	845	35	151	3157
2005–2009	798	35	163	2408
2010–2014	553	35	209	2737
2015–2019	457	63	360	3423
2020–2024	429	49	309	1316
All	3,409	35	210	3423

Table 4: Commissioners with 500 or more recorded votes. Dissent is a vote recorded in the NOES column; the Commission votes by roll call and the minutes record both sides, so this is a census of recorded votes rather than a sample.

Commissioner	Votes	Noes	Dissent rate	Years
Moore	7,608	550	7.2%	2006–2026
Antonini	7,012	264	3.8%	2002–2016
Olague	3,975	309	7.8%	2004–2012
Sugaya	3,734	295	7.9%	2006–2014
S. Lee	3,147	80	2.5%	2002–2008
Koppel	3,099	43	1.4%	2016–2024
W. Lee	2,910	49	1.7%	2002–2008
Fong	2,665	48	1.8%	2010–2019
Joe	2,560	145	5.7%	1998–2002
Chinchilla	2,451	81	3.3%	1998–2002
Borden	2,426	30	1.2%	2008–2014
Theoharis	2,403	48	2.0%	1998–2002
Hillis	2,224	44	2.0%	2012–2019
Hughes	2,175	56	2.6%	2002–2006
Miguel	1,964	32	1.6%	2008–2012
Imperial	1,963	141	7.2%	2020–2026
Diamond	1,850	21	1.1%	2019–2024
Johnson	1,829	28	1.5%	2014–2020
Richards	1,613	51	3.2%	2014–2019
Mills	1,611	20	1.2%	1998–2000

Table 5: Planning Code sections cited in the request text, with the two request types that cite them most. The mapping recovers the Code’s own structure without being told it, which is a validity check on the extraction as much as a finding.

Section	Items	Cited mainly by
§31.04	2,899	<code>conditional_use</code> (1,652), <code>discretionary_review</code> (858)
§303	2,701	<code>conditional_use</code> (2,468), <code>conditional_use_modification</code> (141)
§101.1	923	<code>planning_code_amendment</code> (707), <code>general_plan_amendment</code> (90)
§134	912	<code>variance</code> (519), <code>conditional_use</code> (190)
§302	726	<code>planning_code_amendment</code> (586), <code>rezoning_map_amendment</code> (103)
§317	660	<code>conditional_use</code> (397), <code>discretionary_review</code> (215)
§309	454	<code>downtown_project</code> (252), <code>large_project_authorization</code> (89)
§140	358	<code>variance</code> (166), <code>conditional_use</code> (81)
§209.1	307	<code>conditional_use</code> (288), <code>planning_code_amendment</code> (10)
§321	252	<code>office_allocation</code> (209), <code>large_project_authorization</code> (13)
§209.3	247	<code>conditional_use</code> (227), <code>conditional_use_modification</code> (10)
§304	231	<code>conditional_use</code> (191), <code>conditional_use_modification</code> (21)
§135	214	<code>variance</code> (86), <code>conditional_use</code> (46)
§151	202	<code>variance</code> (106), <code>conditional_use</code> (61)
§329	189	<code>large_project_authorization</code> (179), <code>planning_code_amendment</code> (7)
§121.2	188	<code>conditional_use</code> (167), <code>planning_code_amendment</code> (16)
§161	181	<code>conditional_use</code> (140), <code>downtown_project</code> (13)
§121.1	180	<code>conditional_use</code> (168), <code>planning_code_amendment</code> (7)

Table 6: What the item table can be joined to, and how far it reaches.

Join key	Coverage	Reaches
Assessor block	84.5%	2,557 distinct blocks
Block + lot (APN)	84.1%	8,295 distinct parcels
Building-permit number	22.1%	2,295 distinct permits, 97.6% matched in DBI